

Vector Clocks and Causal Ordering

1. Introduction

In distributed systems, ensuring consistency among different nodes is a challenging task, especially when events occur asynchronously. Traditional logical clocks like Lamport timestamps fail to capture causal dependencies accurately. This project explores and implements **Vector Clocks** to maintain **causal consistency** across a multi-node key-value store. The system simulates a real-world distributed environment using Python, Flask, and Docker, focusing on ensuring that updates respect the causal ordering of events.

2. Objective

To design and implement a distributed key-value store with three or more nodes that:

- Maintains causal consistency using vector clocks
- Buffers and delays writes until causal dependencies are satisfied
- Demonstrates causal ordering through controlled test scenarios
- Is fully containerized using Docker and orchestrated via Docker Compose

3. Technologies Used

- **Language:** Python 3
- **Web Framework:** Flask
- **Containerization:** Docker, Docker Compose
- **IDE:** IntelliJ IDEA (Python Plugin)
- **Version Control:** Git + GitHub

4. System Architecture

The system consists of the following components:

a. Nodes (`node.py`) Each node:

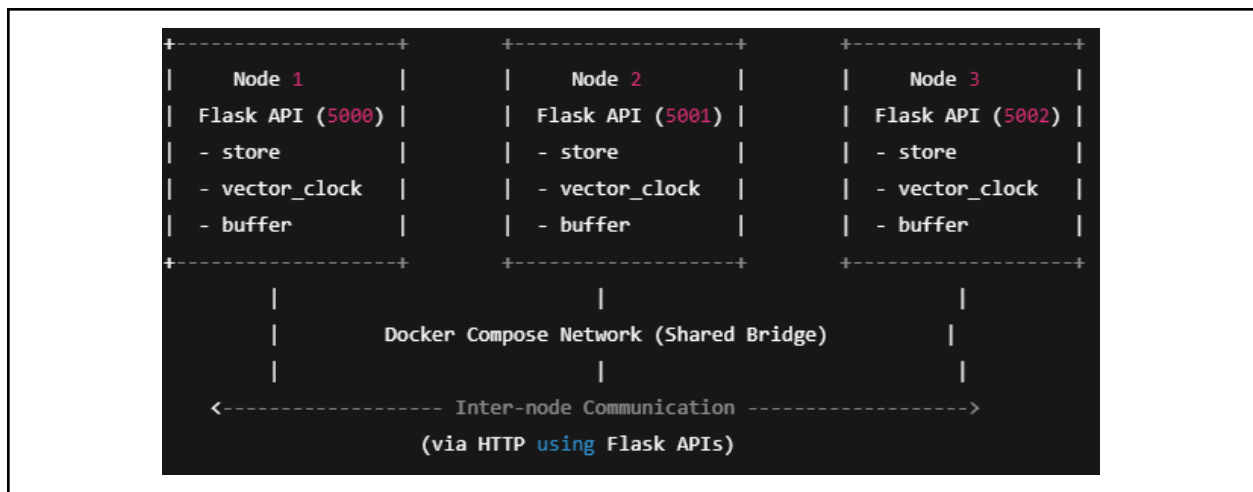
- Runs a Flask server
- Maintains a local key-value store
- Maintains a vector clock with entries for all nodes
- Buffers out-of-order writes

b. Client (`client.py`) Simulates PUT and GET operations to different nodes, mimicking a real user's interaction with the system. This helps demonstrate how causal consistency is enforced.

c. Docker Environment

- Each node runs in its own container
- Docker Compose is used to configure and network the containers

Architecture as follows:



5. Working Mechanism

a. PUT Request (User Write)

- The initiating node increments its vector clock
- Updates its local key-value store
- Sends the update to all peers via a `/write` endpoint

b. WRITE Request (Replication)

- Each peer checks if the incoming write is causally ready (i.e., no missing dependencies)

- If ready, it applies the update
- If not, it buffers the write

c. GET Request

- The node returns the current value of the key and its vector clock

d. Buffer Processing

- After each successful write, the buffer is rechecked for deliverable messages

6. Causal Consistency Logic

Vector clocks are used to track causality. For a message to be causally ready:

```
for node in incoming_vc:
    if incoming_vc[node] > local_vc.get(node, 0):
        return False
```

Only when all conditions are met is the message applied.

7. Test Scenario

Executed using `client.py`:

```
put("node1", "x", "A")
get("node2", "x")
put("node2", "x", "B")
get("node3", "x")
```

This tests whether a write ("B") from node2 is applied only after it has seen the write ("A") from node1.

Expected Output:

- Node2 will buffer "B" until it sees "A"
- Node3 should return the latest causally valid value

8. Challenges Faced

- **Container communication:** Adjusted networking using `host.docker.internal`
- **File path issues in Docker:** Resolved by fixing working directory and volume mounts
- **Causal delivery bugs:** Debugged vector clock mismatch and buffer conditions

9. GitHub Repository

All code is version-controlled and publicly available:

[Vector Clocks and Causal Ordering](#)

10. Conclusion

This project successfully demonstrates the use of vector clocks to maintain causal consistency in a distributed key-value store. The architecture mimics real-world distributed systems, ensuring that updates are applied only when their causal dependencies are met. The implementation is scalable, containerized, and robust against out-of-order message delivery.